

## Lab 3: Cache Simulation

---

### Introduction

In this lab, you will be simulating cache hit ratios or hit rates for direct mapping, N-way associative, and fully associative with LRU and MRU replacement policies. We will be using Google Colaboratory environment (Colab) which is a Jupyter Notebook-like environment where the default coding language is Python, but supports Linux system command line interface and scripting. So instead of installing Linux distribution on virtual machine or installing Python on your computer, Colab provides a user ready interface to readily start our simulation. Also, this lab will also refer you to an online page to help you review some of the concepts for your understanding of cache simulation. After completing this lab, you are required to do the designated homework.

### 1. Concepts from Online Page

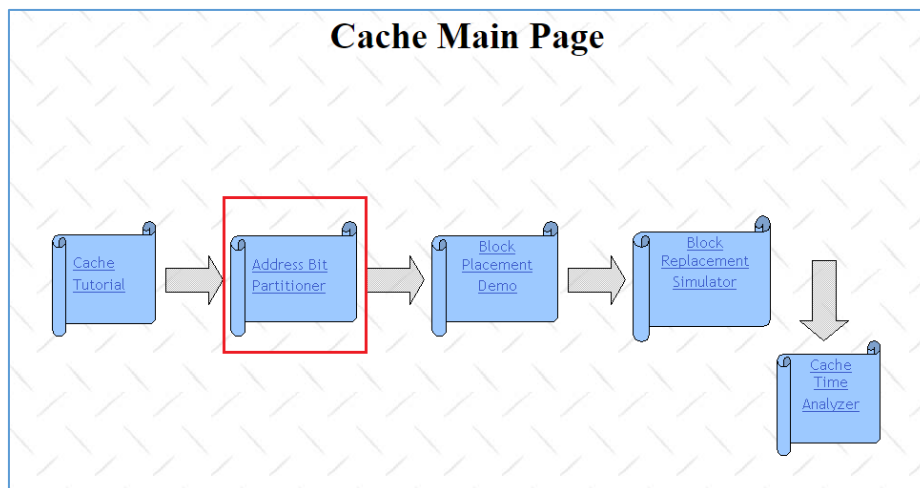
#### Objectives

The objective in this section is just to help you review and visualize some concepts for cache simulation using an online page.

#### A. Cache Address Structure and Placement Visualization

>> Go to this page: <http://www.ecs.umass.edu/ece/koren/architecture/Cache/frame0.htm>

1. Go to the Address Bit Partitioner section.



This part will help you visualize the cache address structure for direct mapping and N-way associative.

Input the settings for **direct mapping**:

**Cache Address Structure**

**Memory Cache Parameters**

Memory Size: 256KB  
 Cache Size: 64KB  
 Block Size: 4B

**Cache Scheme**

Direct Mapping ☒  
 Set Associative ☐  
 Set Size: 4 blocks

Show  
 HELP  
 Return to Main Menu

The output address partition for 256KB main memory, 64KB cache, and 4Byte word alignment (Block Size)

**Address Bit Partitioning**

TAG		INDEX												OFFSET			
17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Compare Bits		Set Select Bits												Byte Select Bits			

The Compare Bits are compared with the corresponding Tag Bits in the Cache Directory.  
 The Set Select Bits are used to select a particular Set in the Cache.  
 The Byte Select Bits are used to select a particular byte in the accessed block.

Memory size = 256KB =  $2^{18}$   
 Block size = 4Bytes =  $2^2$

Number of blocks in cache = Cache size/Block size = 64KB/4B =  $2^{16}/2^2 = 2^{14}$

Number of bits in Tag = Total bits - Index bits - Offset bits =  $18-14-2 = 2$

Settings for **4-way associative mapping**:

**Cache Address Structure**

**Memory Cache Parameters**

Memory Size: 256KB  
 Cache Size: 64KB  
 Block Size: 4B

**Cache Scheme**

Direct Mapping ☐  
 Set Associative ☒  
 Set Size: 4 blocks  
 2 blocks  
 4 blocks  
 8 blocks

Show  
 HELP  
 Return to Main Menu

**Address Bit Partitioning**

TAG		INDEX												OFFSET			
17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Compare Bits		Set Select Bits												Byte Select Bits			

The Compare Bits are compared with the corresponding Tag Bits in the Cache Directory.  
 The Set Select Bits are used to select a particular Set in the Cache.  
 The Byte Select Bits are used to select a particular byte in the accessed block.

Memory size = 256KB =  $2^{18}$   
 Block size = 4Bytes =  $2^2$

Number of sets in cache = Cache size/(Set size \* Block size) = 64KB/(4 blocks \* 4B) =  $2^{16}/(2^2 * 2^2) = 2^{12}$

Number of bits in Tag = Total bits - Index bits - Offset bits =  $18-12-2 = 4$

Settings for **8-way associative mapping**:

**Cache Address Structure**

**Memory Cache Parameters**

Memory Size: 256KB  
 Cache Size: 64KB  
 Block Size: 4B

**Cache Scheme**

Direct Mapping ☐  
 Set Associative ☒  
 Set Size: 8 blocks  
 2 blocks  
 4 blocks  
 8 blocks

Show  
 HELP  
 Return to Main Menu

**Address Bit Partitioning**

TAG		INDEX												OFFSET			
17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Compare Bits		Set Select Bits												Byte Select Bits			

The Compare Bits are compared with the corresponding Tag Bits in the Cache Directory.  
 The Set Select Bits are used to select a particular Set in the Cache.  
 The Byte Select Bits are used to select a particular byte in the accessed block.

Memory size = 256KB =  $2^{18}$   
 Block size = 4Bytes =  $2^2$

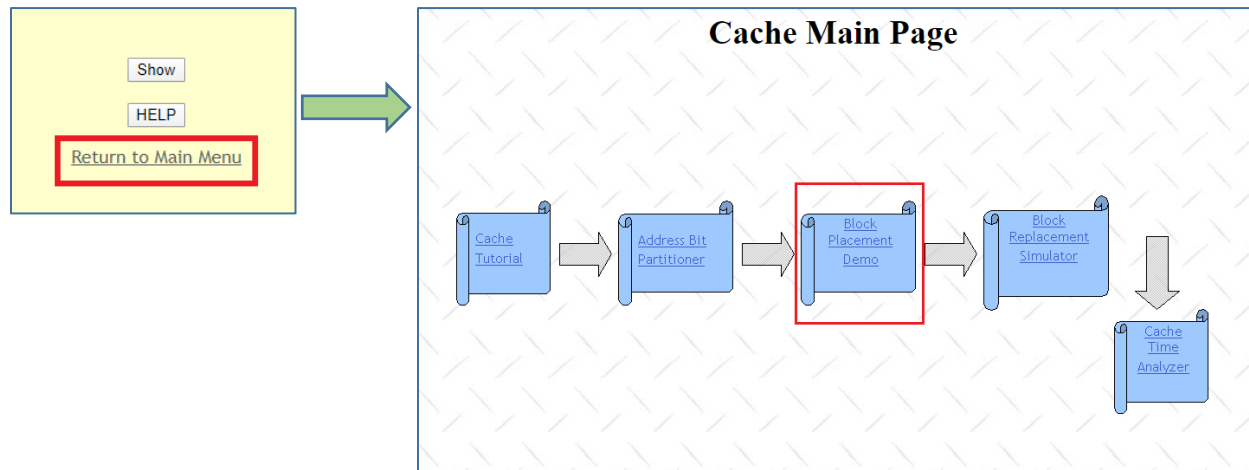
Number of sets in cache = Cache size/(Set size \* Block size) = 64KB/(8 blocks \* 4B) =  $2^{16}/(2^3 * 2^2) = 2^{11}$

Number of bits in Tag = Total bits - Index bits - Offset bits =  $18-11-2 = 5$

Please remember that the INDEX or the set bits are used to map the main memory address to the cache address. So the location of the data in the cache locates to the same set bits (INDEX) in the memory.

Then, the tag bits are compared from the cache location. If the tag bits matches to what the processor is searching then it is a HIT else it is a MISS. The offset bits are mostly ignored for the cache access process.

2. Now, return to the main menu and go to the Block Placement Demo.



This part will help you visualize the content of cache from the given binary address.

>>Input the following Settings:  
  
**Direct Mapping,**  
**128B Cache Size,**  
**4B Block Size,**  
**Binary Address 10111**

**Block Placement for various cache schemes** [Main Menu](#)

Cache Size 128B ▼

Block Size 4B ▼

Direct Mapped Cache ●

Set Associative Cache ☐

10111

☒ Binary

☐ Decimal

MAP

Maximum Decimal Number: 64K

Press MAP.

**OUTPUT**

TAG	INDEX	OFFSET
0 0 0 0	0 0 1 0	1 1

Value in Decimal Format = 5

BLOCK	TAG	INDEX	OFFSET	Value
BLOCK 0				0, 32, 64, 96, 128, 160, ...
BLOCK 1				1, 33, 65, 97, 129, 161, ...
BLOCK 2				2, 34, 66, 98, 130, 162, ...
BLOCK 3				3, 35, 67, 99, 131, 163, ...
BLOCK 4				4, 36, 68, 100, 132, 164, ...
BLOCK 5				5, 37, 69, 101, 133, 165, ...
BLOCK 6				6, 38, 70, 102, 134, 166, ...
BLOCK 7				7, 39, 71, 103, 135, 167, ...
BLOCK 8				8, 40, 72, 104, 136, 168, ...
BLOCK 9				9, 41, 73, 105, 137, 169, ...
BLOCK 10				10, 42, 74, 106, 138, 170, ...
BLOCK 11				11, 43, 75, 107, 139, 171, ...
BLOCK 12				12, 44, 76, 108, 140, 172, ...
BLOCK 13				13, 45, 77, 109, 141, 173, ...
BLOCK 14				14, 46, 78, 110, 142, 174, ...
BLOCK 15				15, 47, 79, 111, 143, 175, ...

>>Now input another address:  
**Binary Address 101001**

101001

☒ Binary

☐ Decimal

MAP

Maximum Decimal Number: 64K



TAG				INDEX					OFFSET	
0	0	0	0	0	1	0	1	0	0	1

Value in Decimal Format = 10

BLOCK 0		0, 32, 64, 96, 128, 160, ...
BLOCK 1		1, 33, 65, 97, 129, 161, ...
BLOCK 2		2, 34, 66, 98, 130, 162, ...
BLOCK 3		3, 35, 67, 99, 131, 163, ...
BLOCK 4		4, 36, 68, 100, 132, 164, ...
BLOCK 5		5, 37, 69, 101, 133, 165, ...
BLOCK 6		6, 38, 70, 102, 134, 166, ...
BLOCK 7		7, 39, 71, 103, 135, 167, ...
BLOCK 8		8, 40, 72, 104, 136, 168, ...
BLOCK 9		9, 41, 73, 105, 137, 169, ...
BLOCK 10		10, 42, 74, 106, 138, 170, ...
BLOCK 11		11, 43, 75, 107, 139, 171, ...
BLOCK 12		12, 44, 76, 108, 140, 172, ...
BLOCK 13		13, 45, 77, 109, 141, 173, ...
BLOCK 14		14, 46, 78, 110, 142, 174, ...
BLOCK 15		15, 47, 79, 111, 143, 175, ...

>>Input the following Settings:

## Set Associative Mapping,

**4-Way,**

### 128B Cache Size,

### 4B Block Size,


### Binary Address 10111

## Direct Mapped Cache

Set Associative Cache ☒

### Associativity

4 way ▼



10111

☒ Binary

☐ Decimal

MAP

Maximum Decimal Number: 64K

>>Now input another address:  
**Binary Address 101001**

101001

☒ Binary

☐ Decimal

MAP


Maximum Decimal Number: 64K

## OUTPUT

TAG		INDEX			OFFSET	
0	0	1	0	1	1	1

Value in Decimal Format = 5

SET 0					0, 8, 16, 24, 32, 40, ...
SET 1					1, 9, 17, 25, 33, 41, ...
SET 2					2, 10, 18, 26, 34, 42, ...
SET 3					3, 11, 19, 27, 35, 43, ...
SET 4					4, 12, 20, 28, 36, 44, ...
SET 5					5, 13, 21, 29, 37, 45, ...
SET 6					6, 14, 22, 30, 38, 46, ...
SET 7					7, 15, 23, 31, 39, 47, ...



## OUTPUT

TAG		INDEX			OFFSET	
0	1	0	1	0	0	1

Value in Decimal Format = 2

SET 0					0, 8, 16, 24, 32, 40, ...
SET 1					1, 9, 17, 25, 33, 41, ...
SET 2					2, 10, 18, 26, 34, 42, ...
SET 3					3, 11, 19, 27, 35, 43, ...
SET 4					4, 12, 20, 28, 36, 44, ...
SET 5					5, 13, 21, 29, 37, 45, ...
SET 6					6, 14, 22, 30, 38, 46, ...
SET 7					7, 15, 23, 31, 39, 47, ...

Notice the change in cache content mapping for direct mapping and 4-way associative mapping with 128B cache size x 4B block size.

>> After introducing this interactive online tool, you can explore it by yourself to review or study cache addressing concepts.

## 2. Cache Simulation in Google Colab

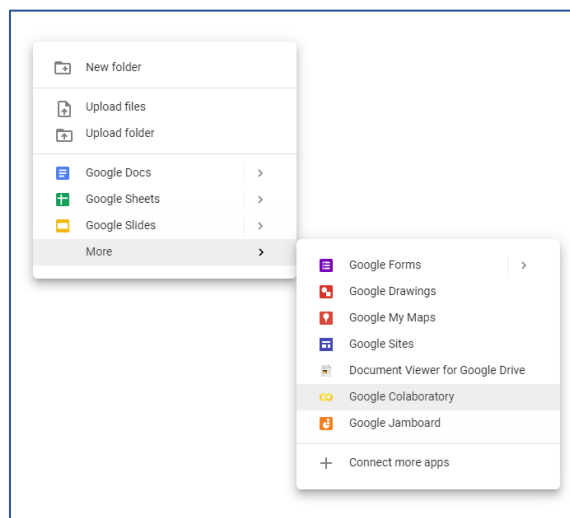
### Objectives

The objective in this section is to introduce the Google Colab environment and simulate different cache configurations.


### A. Starting Google Colab

>> If you don't have a Google account please make one because this is required to access Google Colab.

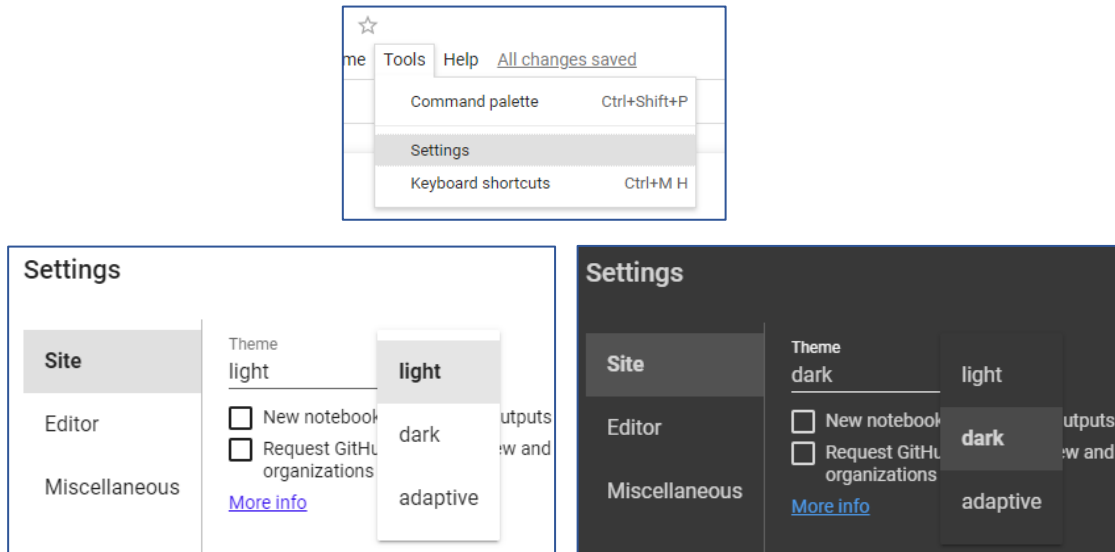
1. Open Google Drive (sign in using your Google account if required or sign in your Chrome browser using your Google account). Then go to your personal drive and right click on any black space then choose Google Colaboratory.



2. To rename your Colab notebook, click the .ipynb filename at the top left part. Then rename it with "lab3\_cachesim" as indicated in the figure and press Enter.




3. To customize the settings of your notebook, click Tools then Settings. You can customize to light or dark theme whichever suits your eyes and enable line numbers that is important for debugging.



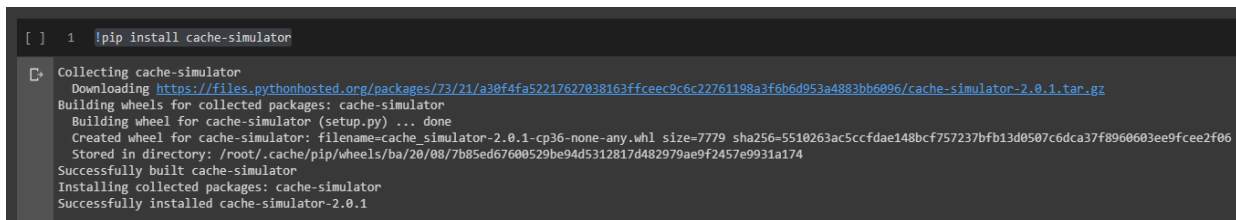
4. Now install the cache-simulator package by inputting the command in the Code Cell. Copy and paste the following command:

```
!pip install cache-simulator
```

Then to run the command press this button beside the cell.  
Or press **Shift + Enter** to run and insert new cell below.




>>After running the command, the package should be successfully installed.

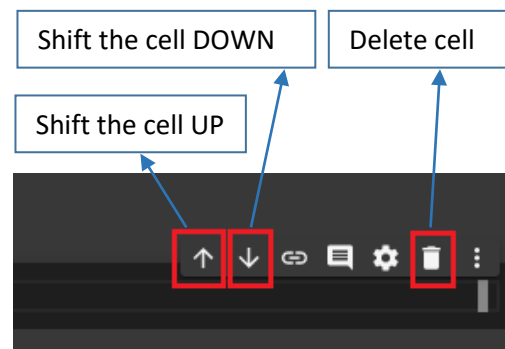


>>Some important functions

To manually insert a cell press the +Code at the upper left of the page.



To show cell options, click the cell and the options will show at the right side of the cell.



## B. Cache Simulation

1. For this package we have to go to its directory to call the function. Copy and paste the command and run the cell.

```
cd '/usr/local/lib/python3.6/dist-packages'
```

>>After running the command, the directory should point to the package directory.

```
[ ] 1 cd '/usr/local/lib/python3.6/dist-packages'
[ ] /usr/local/lib/python3.6/dist-packages
```

2. Now let us test the cache-simulator package with the following configuration.

Cache Size: <b>32W</b>	Blocks per set: <b>1</b>	Block Size: <b>4W</b>	Replacement Policy: <b>LRU</b>	Memory Word Address Size: <b>2<sup>8</sup></b>
Word Addresses: <b>0 4 8 16 4 0 144 16 136 8</b> (each address is separated by a space)				

Copy and paste the following command to simulate the given configuration and run the cell. Notice that this is direct mapping with blocks per set = 1. (DO NOT COPY THE GREEN TEXT)

```
# Direct Mapping
!python cachesimulator --cache-size 32 --num-blocks-per-set 1 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

Code in Colab:

```
[8] 1 # Direct Mapping
2 !python cachesimulator --cache-size 32 --num-blocks-per-set 1 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

Output:

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss		
0	0000 0000	000	000	00	miss		
4	0000 0100	000	001	00	miss		
8	0000 1000	000	010	00	miss		
16	0001 0000	000	100	00	miss		
4	0000 0100	000	001	00	HIT		
0	0000 0000	000	000	00	HIT		
144	1001 0000	100	100	00	miss		
16	0001 0000	000	100	00	miss		
136	1000 1000	100	010	00	miss		
8	0000 1000	000	010	00	miss		
Cache							
000	001	010	011	100	101	110	111
0,1,2,3	4,5,6,7	8,9,10,11	16,17,18,19				

Notice how the second 8 and 16 addresses MISS, this is because addresses 144 and 136 occupies the same cache block as address 8 and 16. **Hit Ratio =  $2 / 10 = 0.2$** . Also, the output will show the final state of the cache on the lower portion.

3. Now, simulate using the same configuration and test addresses as B-2 but **setting block per set = 2**. Notice that this is 2-way associative mapping. Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# 2-Way Associative Mapping
!python cachesimulator --cache-size 32 --num-blocks-per-set 1 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

Code in Colab:

```
[16] 1 # 2-Way Associative Mapping
      2 !python cachesimulator --cache-size 32 --num-blocks-per-set 2 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

Output:

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
0	0000 0000	0000	00	00	miss
4	0000 0100	0000	01	00	miss
8	0000 1000	0000	10	00	miss
16	0001 0000	0001	00	00	miss
4	0000 0100	0000	01	00	HIT
0	0000 0000	0000	00	00	HIT
144	1001 0000	1001	00	00	miss
16	0001 0000	0001	00	00	miss
136	1000 1000	1000	10	00	miss
8	0000 1000	0000	10	00	HIT
Cache					
00	01	10	11		
16,17,18,19	144,145,146,147	4,5,6,7	8,9,10,11	136,137,138,139	

Notice that the second 8 address is a HIT. This is because 2 addresses can now fit using 2-way associative mapping with the given index or set bits. But, the second 16 address is a miss. This is because the replacement policy used is LRU; address 16 is the least recently used after inputting the second address 0. Notice that addresses 0, 16, and 144 have the same index. That means they map to the same set in the cache. This means when address 144 comes, address 16 gets evicted. **Hit Ratio =  $3 / 10 = 0.3$** .

4. Now, simulate using the same configuration and test addresses as B-3 but setting the **replacement policy as MRU**. Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# 2-Way Associative Mapping, MRU replacement policy
!python cachesimulator --cache-size 32 --num-blocks-per-set 2 --num-words-per-block 4 --replacement-policy mru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

Code in Colab:

```
[18] 1 # 2-Way Associative Mapping, MRU replacement policy
      2 !python cachesimulator --cache-size 32 --num-blocks-per-set 2 --num-words-per-block 4 --replacement-policy mru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```



**Output:**

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
0	0000 0000	0000	00	00	miss
4	0000 0100	0000	01	00	miss
8	0000 1000	0000	10	00	miss
16	0001 0000	0001	00	00	miss
4	0000 0100	0000	01	00	HIT
0	0000 0000	0000	00	00	HIT
144	1001 0000	1001	00	00	miss
16	0001 0000	0001	00	00	HIT
136	1000 1000	1000	10	00	miss
8	0000 1000	0000	10	00	HIT
Cache					
00	01	10	11		
144,145,146,147	16,17,18,19	4,5,6,7	8,9,10,11	136,137,138,139	

Using the MRU replacement policy, the second 16 address is now a HIT. Notice that addresses 0, 16, and 144 have the same index 00. With MRU, since the second 0 address hits, when address 144 comes, address 0 will be evicted instead of address 16. **Hit Ratio = 4 / 10 = 0.4.**

5. Now, simulate using the same configuration and test addresses as B-2, but **setting block per set = 4**. Notice that this is 4-way associative mapping. Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# 4-Way Associative Mapping
!python cachesimulator --cache-size 32 --num-blocks-per-set 4 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-
addrs 0 4 8 16 4 0 144 16 136 8
```

**Code in Colab:**

```
[20] 1 # 4-Way Associative Mapping
      2 !python cachesimulator --cache-size 32 --num-blocks-per-set 4 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-
addrs 0 4 8 16 4 0 144 16 136 8
```

**Output:**

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
0	0000 0000	00000	0	00	miss
4	0000 0100	00000	1	00	miss
8	0000 1000	00001	0	00	miss
16	0001 0000	00010	0	00	miss
4	0000 0100	00000	1	00	HIT
0	0000 0000	00000	0	00	HIT
144	1001 0000	10010	0	00	miss
16	0001 0000	00010	0	00	HIT
136	1000 1000	10001	0	00	miss
8	0000 1000	00001	0	00	miss
Cache					
0	1				
8,9,10,11	136,137,138,139	16,17,18,19	144,145,146,147	4,5,6,7	

Notice that the second 16 address is a HIT while the second 8 address is a MISS. With 4-way associative mapping, addresses 0, 8, 16, 144, and 136 maps to the same set. Before replacement, the set needs to be

full, so addresses 0, 8, 16, and 144 will occupy set index 0. Since the second addresses 0, and 16 HITS and address 144 just recently occupies the 4<sup>th</sup> block in index 0, address 8 is the recently used block. So when address 136 will come, address 8 will be evicted. Then, the second address 8 will be a MISS but will replace address 0 (LRU block in this state) and the final state of the cache is displayed. **Hit Ratio = 3 / 10 = 0.3.**

6. Now, simulate using the same configuration and test addresses as B-5 but setting the **replacement policy as MRU**. Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# 4-Way Associative Mapping, MRU replacement policy
!python cachesimulator --cache-size 32 --num-blocks-per-set 4 --num-words-per-block 4 --replacement-policy mru --num-addr-bits 8 --word-
addrs 0 4 8 16 4 0 144 16 136 8
```

Code in Colab:

```
[22] 1 # 4-Way Associative Mapping, MRU replacement policy
      2 !python cachesimulator --cache-size 32 --num-blocks-per-set 4 --num-words-per-block 4 --replacement-policy mru --num-addr-bits 8 --word-
addrs 0 4 8 16 4 0 144 16 136 8
```

Output:

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
0	0000 0000	00000	0	00	miss
4	0000 0100	00000	1	00	miss
8	0000 1000	00001	0	00	miss
16	0001 0000	00010	0	00	miss
4	0000 0100	00000	1	00	HIT
0	0000 0000	00000	0	00	HIT
144	1001 0000	10010	0	00	miss
16	0001 0000	00010	0	00	HIT
136	1000 1000	10001	0	00	miss
8	0000 1000	00001	0	00	HIT
Cache					
0			1		
0,1,2,3 8,9,10,11 136,137,138,139 144,145,146,147			4,5,6,7		

Using the MRU replacement policy, the second 8 address is now a HIT. Since addresses 0, 8, 16, 144, and 136 maps to the same set index 0. After the second address 0 hits, address 144 occupies the 4<sup>th</sup> block of the set. Then, address 16 hits. From this state, address 136 needs to replace one of the blocks in set index 0. Because address 16 just recently hit (used), following MRU policy, it will be evicted and address 136 will replace it keeping address 8 in the set. **Hit Ratio = 4 / 10 = 0.4.**

7. Now, simulate using the same configuration and test addresses as B-2, but **setting block per set = 8**. Notice that this is fully associative mapping since **Set number = Cache Size/block size = N = 8**. Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# Fully Associative Mapping
!python cachesimulator --cache-size 32 --num-blocks-per-set 8 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-
addrs 0 4 8 16 4 0 144 16 136 8
```

Code in Colab:

```
[28] 1 # Fully Associative Mapping
      2 !python cachesimulator --cache-size 32 --num-blocks-per-set 8 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 8 --word-
addrs 0 4 8 16 4 0 144 16 136 8
```

## Output:

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
0	0000 0000	000 000	n/a	00	miss
4	0000 0100	000 001	n/a	00	miss
8	0000 1000	000 010	n/a	00	miss
16	0001 0000	000 100	n/a	00	miss
4	0000 0100	000 001	n/a	00	HIT
0	0000 0000	000 000	n/a	00	HIT
144	1001 0000	100 100	n/a	00	miss
16	0001 0000	000 100	n/a	00	HIT
136	1000 1000	100 010	n/a	00	miss
8	0000 1000	000 010	n/a	00	HIT
Cache					
0,1,2,3 4,5,6,7 8,9,10,11 16,17,18,19 144,145,146,147 136,137,138,139					

Now, let's change the **replacement policy to MRU** and simulate again the fully associative mapped cache. Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# Fully Associative Mapping, MRU replacement policy
!python cachesimulator --cache-size 32 --num-blocks-per-set 8 --num-words-per-block 4 --replacement-policy mru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

## Code in Colab:

```
[29] 1 # Fully Associative Mapping, MRU replacement policy
      2 !python cachesimulator --cache-size 32 --num-blocks-per-set 8 --num-words-per-block 4 --replacement-policy mru --num-addr-bits 8 --word-addr 0 4 8 16 4 0 144 16 136 8
```

## Output:

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
0	0000 0000	000 000	n/a	00	miss
4	0000 0100	000 001	n/a	00	miss
8	0000 1000	000 010	n/a	00	miss
16	0001 0000	000 100	n/a	00	miss
4	0000 0100	000 001	n/a	00	HIT
0	0000 0000	000 000	n/a	00	HIT
144	1001 0000	100 100	n/a	00	miss
16	0001 0000	000 100	n/a	00	HIT
136	1000 1000	100 010	n/a	00	miss
8	0000 1000	000 010	n/a	00	HIT
Cache					
0,1,2,3 4,5,6,7 8,9,10,11 16,17,18,19 144,145,146,147 136,137,138,139					

In fully, associative mapping, all blocks in the cache belongs to the same set. Since there are only 6 distinct addresses mapped to 8 cache blocks. All of the 6 distinct addresses will be mapped distinctively in the cache. So for LRU and MRU replacement policies, the same final cache state and hit ratio will be observed for the given test addresses. **Hit Ratio = 4 / 10 = 0.4.**

8. Now let's simulate the following with a different configuration.

Cache Size: <b>128W</b>	Blocks per set: <b>1</b>	Block Size: <b>4W</b>	Replacement Policy: <b>LRU</b>	Memory Word Address Size: <b>2<sup>11</sup></b>
Word Addresses: <b>52 56 36 40 44 48 72 76 144 148 152 156 160 164 168 172 176 180 184 220 224 228 184 188 144 148 152 156 160 164 168 172 52 56 36 40 44 48 72 76</b> (each address is separated by a space)				

\*The test addresses are taken from the address sequence of executing the bubble\_sort.asm

From the Placement Visualization, this is how the address partitioning looks like:

TAG				INDEX				OFFSET			
0	0	0	0	0	0	0	0	0	0	0	0

**In the online page simulator:**

The cache is byte addressable:  $1W = 4B$   
 $(4B \text{ per } W) \times 4W = 16 \text{ offset address,}$   
 Then  $\log_2 16 = 4 \text{ bits offset}$

This is how the cache abstraction looks like:

BLOCK 0		0, 32, 64, 96, 128, 160, ...
BLOCK 1		1, 33, 65, 97, 129, 161, ...
BLOCK 2		2, 34, 66, 98, 130, 162, ...
BLOCK 3		3, 35, 67, 99, 131, 163, ...
BLOCK 4		4, 36, 68, 100, 132, 164, ...
BLOCK 5		5, 37, 69, 101, 133, 165, ...
BLOCK 6		6, 38, 70, 102, 134, 166, ...
BLOCK 7		7, 39, 71, 103, 135, 167, ...
BLOCK 8		8, 40, 72, 104, 136, 168, ...
BLOCK 9		9, 41, 73, 105, 137, 169, ...
BLOCK 10		10, 42, 74, 106, 138, 170, ...
BLOCK 11		11, 43, 75, 107, 139, 171, ...
BLOCK 12		12, 44, 76, 108, 140, 172, ...
BLOCK 13		13, 45, 77, 109, 141, 173, ...
BLOCK 14		14, 46, 78, 110, 142, 174, ...
BLOCK 15		15, 47, 79, 111, 143, 175, ...
BLOCK 16		16, 48, 80, 112, 144, 176, ...
BLOCK 17		17, 49, 81, 113, 145, 177, ...
BLOCK 18		18, 50, 82, 114, 146, 178, ...
BLOCK 19		19, 51, 83, 115, 147, 179, ...
BLOCK 20		20, 52, 84, 116, 148, 180, ...
BLOCK 21		21, 53, 85, 117, 149, 181, ...
BLOCK 22		22, 54, 86, 118, 150, 182, ...
BLOCK 23		23, 55, 87, 119, 151, 183, ...
BLOCK 24		24, 56, 88, 120, 152, 184, ...
BLOCK 25		25, 57, 89, 121, 153, 185, ...
BLOCK 26		26, 58, 90, 122, 154, 186, ...
BLOCK 27		27, 59, 91, 123, 155, 187, ...
BLOCK 28		28, 60, 92, 124, 156, 188, ...
BLOCK 29		29, 61, 93, 125, 157, 189, ...
BLOCK 30		30, 62, 94, 126, 158, 190, ...
BLOCK 31		31, 63, 95, 127, 159, 191, ...

The cache has 32 blocks:  
 $(128W / 4W = 32)$

Copy and paste the following command and run the cell. **(DO NOT COPY THE GREEN TEXT)**

```
# Direct Mapping
!python cachesimulator --cache-size 128 --num-blocks-per-set 1 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 11 --word-addr 52 56 36 40 44 48 72 76 144 148 152 156 160 164 168 172 176 180 184 220 224 228 184 188 144 148 152 156 160 164 168 172 52 56 36 40 44 48 72 76 6
```

Code in Colab:

```
[39] 1 # Direct Mapping
      2 !python cachesimulator --cache-size 128 --num-blocks-per-set 1 --num-words-per-block 4 --replacement-policy lru --num-addr-bits 11 --word-addr 52 56 36 40 44 48 72 76 144 148 152 156
```

(The command is very long so the side scroll bar will be used to navigate the code in the cell)

Output:

WordAddr	BinAddr	Tag	Index	Offset	Hit/Miss
5200000	110 100	0000	01101	00	miss
5600000	111 000	0000	01110	00	miss
3600000	100 100	0000	01001	00	miss
4000000	101 000	0000	01010	00	miss
4400000	101 100	0000	01011	00	miss
4800000	110 000	0000	01100	00	miss
7200001	001 000	0000	10010	00	miss
7600001	001 100	0000	10011	00	miss
14400010	010 000	0001	00100	00	miss
14800010	010 100	0001	00101	00	miss
15200010	011 000	0001	00110	00	miss
15600010	011 100	0001	00111	00	miss
16000010	100 000	0001	01000	00	miss
16400010	100 100	0001	01001	00	miss
16800010	101 000	0001	01010	00	miss
17200010	101 100	0001	01011	00	miss
17600010	110 000	0001	01100	00	miss
18000010	110 100	0001	01101	00	miss
18400010	111 000	0001	01110	00	miss
22000011	011 100	0001	10111	00	miss
22400011	100 000	0001	11000	00	miss
22800011	100 100	0001	11001	00	miss
18400010	111 000	0001	01110	00	HIT
18800010	111 100	0001	01111	00	miss
14400010	010 000	0001	00100	00	HIT
14800010	010 100	0001	00101	00	HIT
15200010	011 000	0001	00110	00	HIT
15600010	011 100	0001	00111	00	HIT
16000010	100 000	0001	01000	00	HIT
16400010	100 100	0001	01001	00	HIT
16800010	101 000	0001	01010	00	HIT
17200010	101 100	0001	01011	00	HIT
5200000	110 100	0000	01101	00	miss
5600000	111 000	0000	01110	00	miss
3600000	100 100	0000	01001	00	miss
4000000	101 000	0000	01010	00	miss
4400000	101 100	0000	01011	00	miss
4800000	110 000	0000	01100	00	miss
7200001	001 000	0000	10010	00	HIT
7600001	001 100	0000	10011	00	HIT
Cache					
000000000100010000110010000101001100011101000010010101001011011000110101110001100010001					
144,145,146,147148,149,150,151152,153,154,155156,157,158,159160,161,162,16336,37,38					

Hit Ratio = 11 / 40 = 0.275

(Because the memory address size is long, the output address and final state might be very hard to read because it is all cramped up in a confined spacing.)

**HOMEWORK: SIMULATE N-ASSOCIATIVE AND FULLY ASSOCIATIVE**

Use the following configuration:

Cache Size: <b>128W</b>	Blocks per set: <b>X</b>	Block Size: <b>4W</b>	Replacement Policy: <b>LRU</b>	Memory Word Address Size: <b>2<sup>11</sup></b>
Word Addresses: <b>52 56 36 40 44 48 72 76 144 148 152 156 160 164 168 172 176 180 184 220 224 228 184 188 144 148 152 156 160 164 168 172 52 56 36 40 44 48 72 76</b> (each address is separated by a space)				

Simulate the following:

1. 2-Way Associative
2. 4-Way Associative
3. Fully Associative

Then, present the following for each 1, 2, and 3 items:

**A. Draw the visualization of address partitioning and cache abstraction.** Refer to section B-8. Use this online page: <http://www.ecs.umass.edu/ece/koren/architecture/Cache/frame0.htm> to help you (*no fully associative mapping in the online page*). Please remember to draw the address partitioning based on the given configuration of memory address (11 bits).

**B. Show the simulation command** and also the **output**.

**C. Compute the Hit Ratio** and **compare** it to the **direct mapping hit ratio** shown in section B-8.

**D. Explain the result** of each simulation.